



**POLITECNICO**  
MILANO 1863

# ***Explorative study about FogFlow IoT***

## ***framework***

***Matteo Secco (10529606)***

***(matteo.secco@mail.polimi.it)***

***-Mohammad Rahbari Solout  
(10688365)***

***(mohammad.rahbari@mail.polimi.it)***

***Professor : Elisabetta Di Nitto***

***Teacher assistant : Mehrnoosh  
Askarpour***



# POLITECNICO

MILANO 1863

<b>ABSTRACT</b>	<b>3</b>
<b>INTRODUCTION</b>	<b>4</b>
<b>FOGFLOW</b>	<b>4</b>
<b>SCENARIO</b>	<b>5</b>
<b>IMPLEMENTATION</b>	<b>5</b>
Functions description	6
measure_robot_status_time	6
detect_idle	6
bracelets	6
Deploying on FogFlow	6
Preparing the simulation	6
Running the simulation	7
<b>CONCLUSIONS</b>	<b>7</b>
<b>FUTURE WORK</b>	<b>8</b>
<b>RESOURCES</b>	<b>8</b>
Fog functions	8
Github repository	8
Bash scripts	8
<b>REFERENCES</b>	<b>8</b>



**POLITECNICO**  
MILANO 1863

## **ABSTRACT**

We conducted an explorative study on the FogFlow framework, in order to find out the benefits it brings in the perspective of software development and architecture. To conduct this study, we simulated a scenario of a smart factory where humans and robots work together in the same physical space. We found that FogFlow has several advantages to offer in the field of IoT development: the paradigm of fog computing brings the computation closer to where the world phenomena happens, reducing network latency. The software architecture allows to bypass some limitations that the IoT devices may have on a software level using stateful behaviour, but mitigating the disadvantages of statefulness (wrt RESTfulness) thanks to the locality of the computation and a containerized development. FogFlow emerged to be flexible, allowing changes in the infrastructure at runtime.



**POLITECNICO**  
MILANO 1863

## INTRODUCTION

Collecting massive amounts of data in everyday life poses huge challenges for the user to keep control of his/her data in terms of managing access, sharing and protection. The industrial process requires most of the tasks to be performed locally because of delay and security requirements and structured data to be communicated over the Internet to web services and the cloud. To achieve this task, middleware support is required between the industrial environment and the cloud/web services. In this context, fog is a potential middleware that can be very useful for different industrial scenarios. Fog computing can provide local processing support with acceptable latency to actuators and robots in a manufacturing industry. Additionally, as industrial big data are often unstructured, it can be trimmed and refined by the fog locally, before sending it to the cloud. One of the new platform which applies this kind of computations is FogFlow, which is a distributed data processing framework for IoT platforms to support serverless fog computing with regards to device mobility and system reliability. In the following paragraphs we took a research-based and performance-oriented look at this platform based on scenario-driven. First of all you will see a brief description about FogFlow, afterward our scenario about a Smart Factory for having a real interaction assessment and applying it on the mentioned platform. here are some of its specifications:

## FOGFLOW

FogFlow is an IoT edge computing framework to automatically orchestrate dynamic data processing flows over cloud and edges driven by context, which is the unique feature of this platform. Therefore FogFlow is capable to orchestrate data processing based on three types of contexts:

- System context: available resources which are changing over time.
- Data context: the structure and registered metadata of available data, including both raw sensor data and intermediate data.
- Usage context: high level intents defined by all different types of users (developers, service consumers, data providers) to specify what they want to achieve.

This three contexts empower FogFlow to orchestrate IoT services in a more intelligent and automatic manner. Context orchestration is the key factor to distinguish this platform from the others such as EdgeX, Azure IoT Edge, Amazon Greengrass. This aspect is enabled by the design of introducing a new layer, namely IoT Discovery, which provides a update summary of available entity data on all brokers. As compared to event or topic based orchestration, our context-based orchestration in FogFlow is more flexible and more lightweight. This is because the orchestration decisions in FogFlow can be made based on aggregated context, without reading



**POLITECNICO**  
MILANO 1863

through all involved data streams.  
Here you can see briefly some key  
features of FogFlow:

- A. Context-driven orchestration  
mechanism
- B. Serverless edge computing
- C. Based on Docker
- D. Dynamic data orchestrator



## SCENARIO

PoliBot is a (fictional) smart factory where robots and humans work and move around in the same space. An existing IoT infrastructure exists, to enforce safety policies around the factory. In particular, two safety rules have to be verified:

- A robot must not be in Idle for more than a given threshold. If this happens, it is likely that the robot is experiencing some kind of malfunctioning, and is to be repaired.
- At any moment, robots and humans should be not closer that a fixed distance. This is to ensure that no human may accidentally get harmed by a robot.

The existing infrastructure collects position and status (Idle, Moving, Working) of the robots, and uses smart bracelets to collect the position of the humans.

One day, PoliBot decided to open other factories in different locations. Each factory comes with the same IoT infrastructure, but there is the desire from the direction to have aggregated data about the violation of the rules above available in the central office.

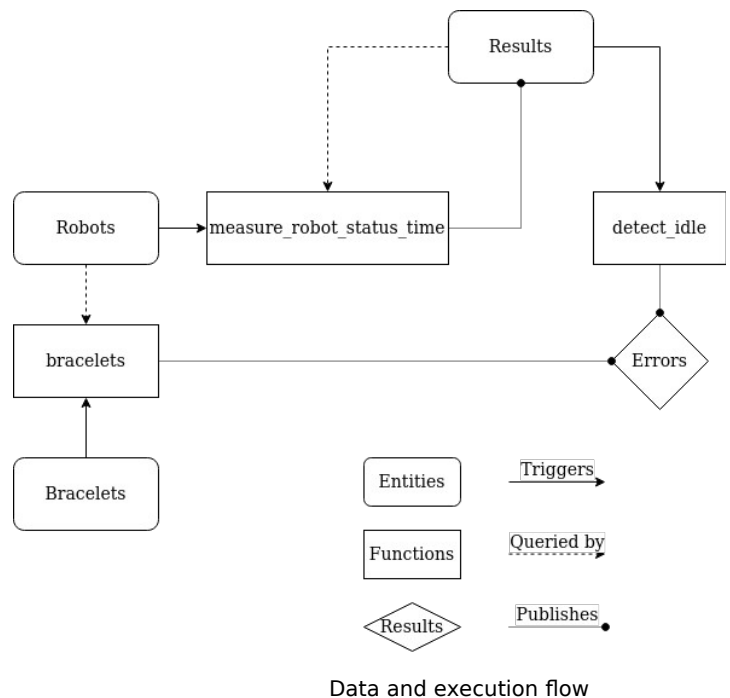
Our goal was to simulate the application of the FogFlow framework to this scenario, in order to study advantages and disadvantages of the framework itself.

## IMPLEMENTATION

At the beginning of our explorative research, we wanted to develop the safety rules described in the [Scenario](#) section. We implemented 3 fog-functions, written in java using the Spring Boot framework, and available on Docker Hub:

- [measure\\_robot\\_status\\_time](#)
- [detect\\_idle](#)
- [bracelets](#)

These functions interact each other and with the environment according to the following diagram:





## POLITECNICO MILANO 1863

### Functions description

Follows a summary of the behaviour of each function:

#### measure\_robot\_status\_time

This function is triggered by any entity update of type Robot. Let's call robot the entity updated. On start, it queries FogFlow for an entity of type Result through the available ones having as id Result.robot-id. Three cases are possible:

- Such an entity doesn't exist: this means that the received update is the first update for robot. → A new Result is published, having result-id=robot-id, result-time=actual time measured by the server, result-last\_interval=0, result-status=robot-status.
- Such a result entity exists, and result-status  $\neq$  robot-status → result is updated, setting result-time=actual time measured by the server, result-last\_interval=0, result-status=robot-status.
- Such a result entity exists, and result-status = robot-status → result is updated, setting result-time=actual time measured by the server, result-status=robot-status, result-last\_interval=result-last\_interval + actual time - result-time.

#### detect\_idle

This function is triggered by an entity update of type Result (result). If result-last\_interval is greater than a fixed

threshold, the function publishes an Error entity indicating the Id of the robot that was idle for too long (the Id is extracted from the result Id), how long it has been in idle, and at which time the error was detected.

#### bracelets

This function is triggered by an entity update of type Bracelet (let bracelet be the entity that triggered the update). It queries FogFlow for all the available Robot entities. Then, for each of them, calculates the distance between the robot and bracelet. If the resulting distance is greater than a fixed threshold, an error is published indicating the Ids of both the robot and bracelet, and the positions of both

### Deploying on FogFlow

When the code was ready, we registered the functions into the fogflow system. This operation consists of three steps:

1. Register an operator for each function
2. Link a docker image to each operator (more than one image can be linked to each operator if we want to provide different code for different operating systems or hardware architectures, but this wasn't our case)
3. Register a fog function for each operator

At the beginning we used to perform such operations by the FogFlow dashboard, but with time we preferred to automate the process using simple [bash scripts](#) performing curl requests.



## POLITECNICO MILANO 1863

After developing isolated functions, we then tried to structure more the use case developing a service topology responsible of running the functions and managing the data.

### Preparing the simulation

We built a simulation of a factory written in Node.js. We created it editing the [powerpanel simulation](#) available on FogFlow public repository. We started by developing the Robot entity, with movement and working capabilities. The entities work in discrete time, so at each update, the robot may perform a fixed percentage of the current job, move by a fixed distance, change status, or do nothing (in case it is in idle). In any case, a robot will never change status while finishing a job, and when it finishes the job it will perform an attempt to change status.

The percentage completed at each update, the distance travelled, and the probability of changing status, are coded into a json file, named [robot\\_profile.json](#).

After the robot has been completed we prepared the Bracelet entity, which is in fact a robot always in the “moving” status.

### Running the simulation

In order to run the simulation, we first had to setup FogFlow on our machines as explained in the [official documentation](#).

Then, we had to update the IP address of the discoveryURL property contained in the [simulation\\_profile.json](#) file according to the address of the machine where

FogFlow was running.

Finally, we needed to register the operators and fog functions as explained in the section [Deploying on FogFlow](#). In case we used the scripts, we needed to provide the same IP address used before.

After all these steps were done, we were able to run the simulation and see the results on the FogFlow dashboard: the resulting errors will get published as entities and appear in the [System Status/Device](#) section.

### Service Topology

We tried to collect the functions into a service topology. Having the functions ready and unchanged, all we had to do was to specify the data flow. Following the documentation was enough to setup the Topology, but a curious behavior emerged: the Topology was only triggered if the operators it used was assigned to an independent FogFunction, and would not if the operator was only used in the topology. This is because the Service Topology uses the FogFunctions to instantiate the tasks, as can be proved by comparing the ID of the running instances in Service Topology>TaskInstance with those in Fog Function>TaskInstance by the dashboard: the running IDs are the same

## EDGE COMPUTING

In the last phase of our study, we wanted to emulate a real distributed system to test how FogFlow behaved with regard to edge computing. In order to do so, we started 1 cloud node and declared its location in Lodi,





## POLITECNICO MILANO 1863

Italy; and 2 edge nodes in Milan and Genoa, Italy. The actual nodes were running on a single machine and were communicating by virtual interfaces.

As expected, updates regarding changes in the existing nodes, registered Operators, Fog Functions or Service Topologies are immediately propagated to all the nodes. Also the information about running nodes (status and location) were shared between all the nodes, and also the information about which node was running instances of some function.

The assignment of brokers to clients also didn't as expected: querying multiple times for nearby brokers providing the same client location gave inconsistent results. So we tried moving the third node from Genoa to Beijing, China. This didn't change the result, it seems that the available brokers are just alternated at each request following a round robin ordering.

The situation is similar with the execution of Fog Functions: in this case, in our simulation no matter where the data came from, they were all processed in the node of Milan.

### CONCLUSIONS

FogFlow emerged in the beginning as an interesting research prototype. The containerized structure adopted for the fog functions brings all the advantages of containerization: portability, agility, speed, fault isolation, efficiency, ease of management.

FogFlow is language-agnostic, and any component of a development team

could code in any programming language he want: all the data are shared by the NGSI protocol.

Compared to other serverless technologies, in particular with [AWS Lambda](#), FogFlow exhibit some peculiarities: it is distributed by design, allowing the developers to bring the computation closer to where the world phenomena happens (so to get improvements in network latency), and to decentralize the information (bringing both performance and security improvements).

The framework is stateful, peculiarity that allows to bypass eventual limitations of the IoT devices by software (like we did to track for how long a robot remained in a state). This could be a good practice for the purpose of the framework: it reduces the costs of the IoT devices needed, and the tradeoff (losing the advantages of a RESTful architecture) is outbalanced by the locality of computation and containerized paradigm.

The infrastructure can be changed at runtime (adding cloud and edge nodes). The change is not perfect yet (for example, when adding a node the existing devices are not moved to that node, even if it is more convenient), but works perfectly fine in the perspective of this scenario (in particular, assuming to have sticky clients). Anyway this may be an issue in the perspective of more complex scenarios, for example a smart city.

Testing FogFlow in the optic of edge computing was indeed disappointing: the brokers do not take geographical distance into account when assigning



# POLITECNICO MILANO 1863

the client to brokers or workers, and so the system doesn't behave like an edge computing framework at all.

## FUTURE WORK

There are many features of FogFlow that we didn't explore but may be interesting for the scenario we analyzed. In particular, we would like to explore more about the integration with Grafana, that would allow a user-friendly data representation for the Business people and experimenting with the entity aggregation mechanisms offered by FogFlow.

## RESOURCES

### Fog functions

- [measure\\_robot\\_status\\_time](#)
- [detect\\_idle](#)
- [bracelets](#)

### Github repository

- [smart\\_factory](#)

### Bash scripts

- [register\\_operator](#)
- [register\\_fog\\_function](#)

## REFERENCES

- B. Cheng, J. Fürst, G. Solmaz, T. Sanada, "[Fog Function: Serverless Fog Computing for Data Intensive IoT Services](#)," in the proceedings of 2019 IEEE Conference on Service Computing (IEEE SCC'19) (Won the best paper award), Milan, 2019, pp.28-35

- M. Fadel Argerich, B. Cheng, J. Fuerst, "[Reinforcement Learning based Orchestration for Elastic Services](#)", in the proceedings of the 5th IEEE World Forum on Internet of Things, WF-IoT 2019, Limerick, Ireland, April 15-18, 2019, pp. 352-357
- B. Cheng, E. Kovacs, A. Kitazawa, K. Terasawa, T. Hada, M. Takeuchi, "[FogFlow: Orchestrating IoT Services over Cloud and Edges](#)", NEC Technical Journal, 2018/11
- B. Cheng, G. Solmaz, F. Cirillo, E. Kovacs, K. Terasawa and A. Kitazawa, "[FogFlow: Easy Programming of IoT Services Over Cloud and Edges for Smart Cities](#)", in IEEE Internet of Things Journal, 2017
- [FogFlow official documentation](#)